

# project-1

February 21, 2024

```
[ ]: LOGISTIC REGRESSION BASED ON PROJECTS
```

```
[ ]: Students Performance based on Logistic Regression
```

```
[96]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score
      from sklearn.model_selection import train_test_split
```

```
[5]: import pandas as pd
```

```
[97]: df=pd.read_csv(r"C:\Users\micro\Downloads\LR_Student_Performance.csv")
      df
```

```
[97]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	Yes	9	
1	4	82	No	4	
2	8	51	Yes	7	
3	5	52	Yes	5	
4	7	75	No	8	
...	...	...	...	...	
9995	1	49	Yes	4	
9996	7	64	Yes	8	
9997	6	83	Yes	8	
9998	9	97	Yes	7	
9999	7	74	No	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0
...	...	...
9995	2	23.0
9996	5	58.0
9997	5	74.0
9998	0	95.0

9999

1

64.0

[10000 rows x 6 columns]

```
[40]: x=df[['Hours Studied']]
      y=df['Previous Scores']
      from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪4,random_state=100)
      from sklearn.linear_model import LogisticRegression
      model=LogisticRegression()
      model
```

```
[40]: LogisticRegression()
```

```
[41]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪2,random_state=101)
```

```
[42]: model.fit(x_train,y_train)
```

```
[42]: LogisticRegression()
```

```
[43]: y_pred=model.predict(x_test)
      y_pred
```

```
[43]: array([54, 54, 86, ..., 54, 54, 86], dtype=int64)
```

```
[44]: y_test
```

```
[44]: 6676    55
      6421    98
      9834    57
      8492    71
      9982    51
      ..
      4441    65
      4166    65
      2567    62
      8527    64
      406     42
      Name: Previous Scores, Length: 2000, dtype: int64
```

```
[45]: from sklearn.metrics import accuracy_score
```

```
[47]: import numpy as np
```

```
[98]: acc=accuracy_score(y_test,np.round(y_pred))
      acc
```

```
[98]: 0.015
```

```
[99]: inputdata=[[17]]
      prediction=model.predict(inputdata)
      prediction
```


```
C:\Users\micro\AppData\Local\Programs\Python\Python312\Lib\site-
packages\sklearn\base.py:493: UserWarning: X does not have valid feature names,
but LogisticRegression was fitted with feature names
  warnings.warn(
```

```
[99]: array([56], dtype=int64)
```

```
[100]: from sklearn.metrics import mean_squared_error
      mse=mean_squared_error(y_test,y_pred)
      mse
```

```
[100]: 557.3895
```

```
[70]: import seaborn as sns
      import matplotlib.pyplot as plt
      sns.lmplot(x="Sleep Hours",y="Sample Question Papers_
      ↳Practiced",data=df,scatter_kws={"color":'violet' },line_kws={'color':"red"})
      sns.set_style('whitegrid')
      ax=plt.gca()
      plt.gca()
      plt.gca().set_facecolor('white')
```



```
[55]: #To check duplicate values
duplicate_rows=df.duplicated()
df[duplicate_rows].sum()
```

```
[55]: Hours Studied
642
Previous Scores
8865
Extracurricular Activities
NoYesNoYesYesNoNoNoNoYesNoYesNoNoNoYesYesYes...
Sleep Hours
819
Sample Question Papers Practiced
585
Performance Index
7094.0
dtype: object
```

```
[57]: x=df.drop("Hours Studied",axis=1)
x
```

```
[57]:
```

	Previous Scores	Extracurricular Activities	Sleep Hours \
0	99	Yes	9
1	82	No	4
2	51	Yes	7
3	52	Yes	5
4	75	No	8
...	...	...	...
9995	49	Yes	4
9996	64	Yes	8
9997	83	Yes	8
9998	97	Yes	7
9999	74	No	8

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0
...	...	...
9995	2	23.0
9996	5	58.0
9997	5	74.0
9998	0	95.0
9999	1	64.0

[10000 rows x 5 columns]

```
[58]: x.shape
```

```
[58]: (10000, 5)
```

```
[101]: y=df.drop("Sleep Hours",axis=1)
y
```

```
[101]:
```

	Hours Studied	Previous Scores	Extracurricular Activities \
0	7	99	Yes
1	4	82	No
2	8	51	Yes
3	5	52	Yes
4	7	75	No
...	...	...	...
9995	1	49	Yes
9996	7	64	Yes

9997	6	83	Yes
9998	9	97	Yes
9999	7	74	No

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0
...	...	...
9995	2	23.0
9996	5	58.0
9997	5	74.0
9998	0	95.0
9999	1	64.0

[10000 rows x 5 columns]

```
[64]: print("Before dropping duplicate:",df.shape)
      df.drop_duplicates()
      print("After dropping duplicate:",df.shape)
```

Before dropping duplicate: (10000, 6)

After dropping duplicate: (10000, 6)

```
[66]: response=df["Extracurricular Activities"]
      response.dtype
```


```
[66]: dtype('O')
```

```
[68]: response=df["Performance Index"]
      response.dtype
```

```
[68]: dtype('float64')
```


```
[75]: sns.jointplot(x=response.index,y="Previous_
      ↳Scores",data=df,kind='hex',color='pink',edgecolor='black')
```

```
[75]: <seaborn.axisgrid.JointGrid at 0x26d970a52e0>
```



```
[76]: sns.violinplot(response)
```

```
[76]: <Axes: ylabel='Performance Index'>
```



```
[91]: ma=(df["Performance Index"]).max()
      ma
```

```
[91]: 100.0
```

```
[92]: mi=(df["Performance Index"]).min()
      mi
```

```
[92]: 10.0
```

```
[93]: (df['Performance Index']==mi).sum()
```

```
[93]: 1
```

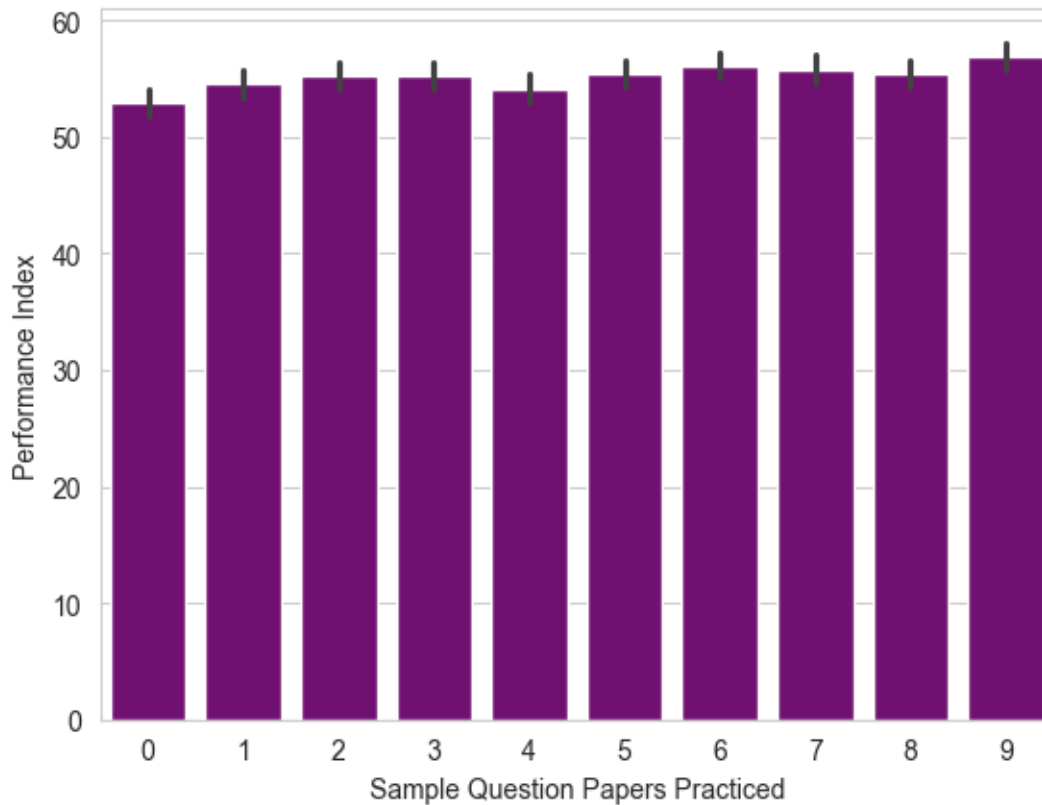
```
[94]: (df['Hours Studied']).unique()
```

```
[94]: array([7, 4, 8, 5, 3, 6, 2, 1, 9], dtype=int64)
```

```
[81]: import seaborn as sns
      sns.barplot(y="Performance Index",x="Sample Question Papers_
      ↪Practiced",data=df,color='purple')
```

```
[81]: <Axes: xlabel='Sample Question Papers Practiced', ylabel='Performance Index'>
```





```
[25]: df.head()
```

```
[25]:
```

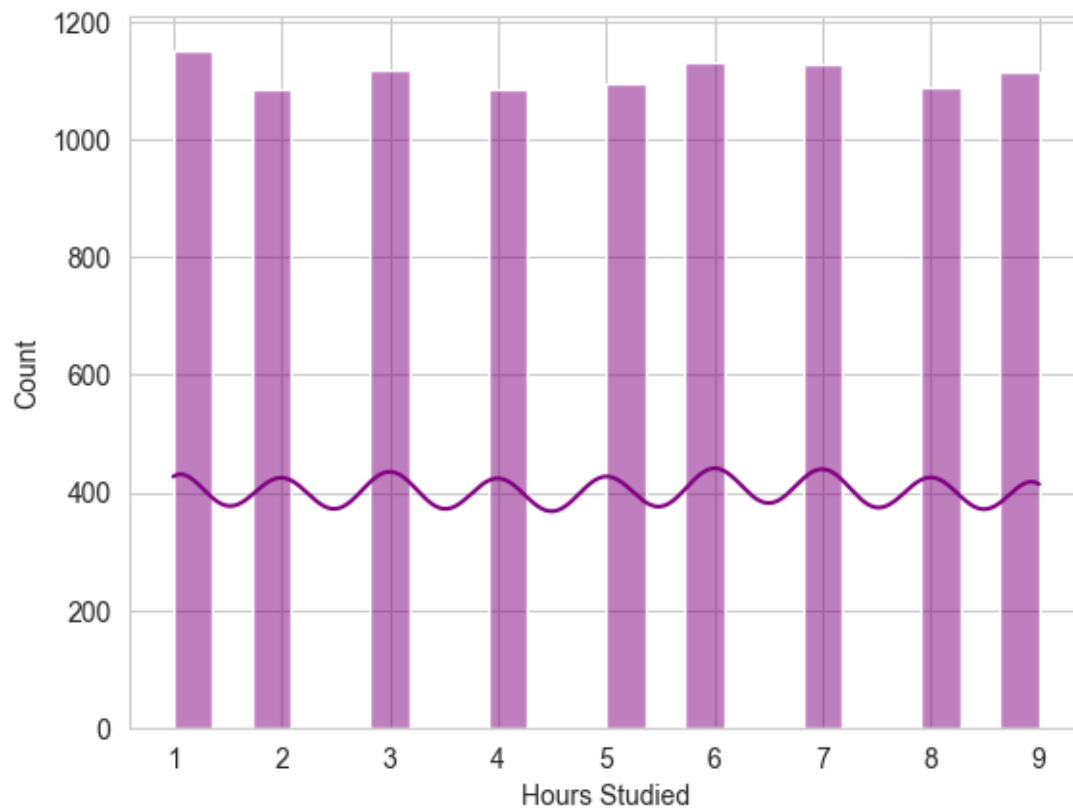
	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	Yes	9	
1	4	82	No	4	
2	8	51	Yes	7	
3	5	52	Yes	5	
4	7	75	No	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0


```
[83]: import seaborn as sns
x=df['Hours Studied']
sns.histplot(x,color='purple',kde=True)
```

```
[83]: <Axes: xlabel='Hours Studied', ylabel='Count'>
```



```
[88]: sns.boxplot(x=df['Hours Studied'],y=df['Performance Index'],color='black')
```

```
[88]: <Axes: xlabel='Hours Studied', ylabel='Performance Index'>
```



```
[89]: x=df[['Hours Studied', 'Previous Scores', 'Extracurricular Activities',
          'Sleep Hours', 'Sample Question Papers Practiced']]
      y=df['Performance Index']
      df.head()
```

```
[89]:   Hours Studied  Previous Scores  Extracurricular Activities  Sleep Hours  \
0              7              99                Yes            9
1              4              82                No             4
2              8              51                Yes            7
3              5              52                Yes            5
4              7              75                No             8

      Sample Question Papers Practiced  Performance Index
0                                1          91.0
1                                2          65.0
2                                2          45.0
3                                2          36.0
4                                5          66.0
```

```
[95]: df['Extracurricular Activities']=df['Extracurricular Activities'].apply(lambda x:
    ↪ x: 1 if x=="Yes"else 0)
```

```
df.head()
```

```
[95]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	1	9	
1	4	82	0	4	
2	8	51	1	7	
3	5	52	1	5	
4	7	75	0	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```
[28]: df.info
```

```
[28]: <bound method DataFrame.info of
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	7	99	Yes	9	
1	4	82	No	4	
2	8	51	Yes	7	
3	5	52	Yes	5	
4	7	75	No	8	
...	...	...	...	...	
9995	1	49	Yes	4	
9996	7	64	Yes	8	
9997	6	83	Yes	8	
9998	9	97	Yes	7	
9999	7	74	No	8	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0
...	...	...
9995	2	23.0
9996	5	58.0
9997	5	74.0
9998	0	95.0
9999	1	64.0

```
[10000 rows x 6 columns]>
```

```
[33]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 6 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   Hours Studied                          10000 non-null  int64
 1   Previous Scores                        10000 non-null  int64
 2   Extracurricular Activities             10000 non-null  object
 3   Sleep Hours                            10000 non-null  int64
 4   Sample Question Papers Practiced       10000 non-null  int64
 5   Performance Index                      10000 non-null  float64
dtypes: float64(1), int64(4), object(1)
memory usage: 468.9+ KB
```

```
[34]: df.isna()
```

```
[34]:
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	\
0	False	False	False	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	
4	False	False	False	False	
...	...	...	...	...	
9995	False	False	False	False	
9996	False	False	False	False	
9997	False	False	False	False	
9998	False	False	False	False	
9999	False	False	False	False	

	Sample Question Papers Practiced	Performance Index
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...	...	...
9995	False	False
9996	False	False
9997	False	False
9998	False	False
9999	False	False

[10000 rows x 6 columns]

```
[36]: df.isna().sum()
```

```
[36]: Hours Studied          0
      Previous Scores        0
      Extracurricular Activities 0
      Sleep Hours            0
      Sample Question Papers Practiced 0
      Performance Index      0
      dtype: int64
```

```
[37]: type(df)
```

```
[37]: pandas.core.frame.DataFrame
```

```
[14]: # To train the algorithm
      clf.fit(x_train,y_train)
```

```
[14]: LogisticRegression()
```

```
[15]: y_pred=clf.predict(x_test)
      y_pred
```

```
[15]: array([0, 0, 0, 2, 1, 2, 1, 1, 2, 0, 2, 0, 0, 2, 2, 1, 1, 1, 0, 2, 1, 0,
        1, 1, 1, 1, 1, 2, 0, 0])
```

```
[ ]:
```