

# project-2

February 21, 2024

```
[ ]: LOGISTIC REGRESSION BASED ON User_Data
```

```
[ ]: Salaries analysis based on logistic regression
```

```
[73]: import pandas as pd
```

```
[74]: data=pd.read_csv(r"C:\Users\micro\Downloads\User_Data.csv")
```

```
[67]: data
```

```
[67]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
..	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

[400 rows x 5 columns]

```
[116]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
```

```
[96]: x=data[['EstimatedSalary']]
y=data['Age']
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪4,random_state=100)
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model
```

```
[96]: LogisticRegression()
```

```
[100]: model.fit(x_train, y_train)
```

```
C:\Users\micro\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[100]: LogisticRegression()
```

```
[98]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.  
↪2,random_state=101)
```

```
[99]: y_pred=model.predict(x_test)  
y_pred
```

```
[99]: array([35, 35, 37, 35, 35, 35, 35, 35, 35, 35, 35, 37, 35, 37, 35, 35, 35, 37,  
35, 35, 35, 37, 37, 35, 35, 35, 35, 35, 35, 35, 35, 35, 35,  
37, 35, 37, 35, 35, 35, 35, 35, 35, 35, 37, 35, 35, 35, 35, 35, 35,  
37, 35, 35, 35, 35, 37, 35, 35, 35, 35, 37, 35, 37, 35, 35, 35, 35,  
35, 35, 35, 35, 37, 35, 35, 35, 35, 35, 35, 35], dtype=int64)
```

```
[101]: y_test
```

```
[101]: 38      26  
387     39  
270     43  
181     31  
195     34  
      ..  
130     31  
13      32  
141     18  
304     40  
167     35  
Name: Age, Length: 80, dtype: int64
```

```
[112]: clf=LogisticRegression()  
clf.fit(x_train,y_train)
```

```
C:\Users\micro\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:  
    https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
    n_iter_i = _check_optimize_result(
```

```
[112]: LogisticRegression()
```

```
[113]: y_pred=clf.predict(x_test)  
y_pred
```

```
[113]: array([35, 35, 36, 35, 35, 35, 35, 35, 35, 35, 35, 36, 37, 37, 35, 35, 35, 36,  
            35, 35, 35, 36, 36, 35, 35, 35, 35, 35, 35, 35, 35, 35, 35, 35,  
            32, 37, 36, 35, 35, 35, 35, 35, 35, 35, 36, 35, 35, 35, 35, 35, 35,  
            37, 35, 35, 35, 35, 36, 35, 35, 35, 35, 37, 35, 37, 35, 35, 35, 35,  
            35, 35, 35, 35, 36, 35, 35, 35, 35, 35, 35, 35], dtype=int64)
```

```
[102]: from sklearn.metrics import accuracy_score
```

```
[103]: import numpy as np
```

```
[110]: accuracy=accuracy_score(y_test,np.round(y_pred))
```

```
[111]: accuracy
```

```
[111]: 0.1
```

```
[117]: inputdata=[[17]]  
prediction=model.predict(inputdata)  
prediction
```

```
C:\Users\micro\AppData\Local\Programs\Python\Python312\Lib\site-  
packages\sklearn\base.py:493: UserWarning: X does not have valid feature names,  
but LogisticRegression was fitted with feature names  
    warnings.warn(
```

```
[117]: array([35], dtype=int64)
```

```
[47]: data.head()
```

```
[47]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
[48]: data.Gender.value_counts()
```

```
[48]: Gender
Female    204
Male      196
Name: count, dtype: int64
```

```
[68]: data=pd.get_dummies(data,columns=["Gender"],dtype=int,drop_first=True)
```

```
[69]: data
```

```
[69]:
```

	User ID	Age	EstimatedSalary	Purchased	Gender_Male
0	15624510	19	19000	0	1
1	15810944	35	20000	0	1
2	15668575	26	43000	0	0
3	15603246	27	57000	0	0
4	15804002	19	76000	0	1
..	...	...	...	...	...
395	15691863	46	41000	1	0
396	15706071	51	23000	1	1
397	15654296	50	20000	1	0
398	15755018	36	33000	0	1
399	15594041	49	36000	1	0

```
[400 rows x 5 columns]
```

```
[50]: (data['Gender_Male']==1).sum()
(data['Gender_Male']==0).sum()
```

```
[50]: 204
```

```
[51]: data
```

```
[51]:
```


	User ID	Age	EstimatedSalary	Purchased	Gender_Male
0	15624510	19	19000	0	1
1	15810944	35	20000	0	1
2	15668575	26	43000	0	0
3	15603246	27	57000	0	0
4	15804002	19	76000	0	1
..	...	...	...	...	...

395	15691863	46	41000	1	0
396	15706071	51	23000	1	1
397	15654296	50	20000	1	0
398	15755018	36	33000	0	1
399	15594041	49	36000	1	0

[400 rows x 5 columns]

```
[52]: import seaborn as sns
sns.barplot(x="Age",y="EstimatedSalary",data=data)
```

```
[52]: <Axes: xlabel='Age', ylabel='EstimatedSalary'>
```



```
[54]: x=data.drop("Purchased",axis=1)
x
```

```
[54]:
```

	User ID	Age	EstimatedSalary	Gender_Male
0	15624510	19	19000	1
1	15810944	35	20000	1
2	15668575	26	43000	0
3	15603246	27	57000	0
4	15804002	19	76000	1

```

..      ... ..
395  15691863  46      41000      0
396  15706071  51      23000      1
397  15654296  50      20000      0
398  15755018  36      33000      1
399  15594041  49      36000      0

```

[400 rows x 4 columns]

```
[55]: x.shape
```

```
[55]: (400, 4)
```

```
[56]: y=data.drop("Gender_Male",axis=1)
y
```

```
[56]:
   User ID  Age  EstimatedSalary  Purchased
0   15624510   19             19000           0
1   15810944   35             20000           0
2   15668575   26             43000           0
3   15603246   27             57000           0
4   15804002   19             76000           0
..      ... ..
395  15691863   46             41000           1
396  15706071   51             23000           1
397  15654296   50             20000           1
398  15755018   36             33000           0
399  15594041   49             36000           1

```

[400 rows x 4 columns]

```
[57]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪2,random_state=101)
```

```
[58]: y.shape
```

```
[58]: (400, 4)
```

```
[59]: x.shape
```

```
[59]: (400, 4)
```

```
[60]: x_train.shape
```

```
[60]: (320, 4)
```

```
[61]: y_train.shape
```

[61]: (320, 4)

```
[62]: from sklearn.linear_model import LogisticRegression
log_model=LogisticRegression()
log_model
```

[62]: LogisticRegression()

```
[71]: data.head()
```

```
[71]:      User ID  Age  EstimatedSalary  Purchased  Gender_Male
0    15624510   19         19000         0         1
1    15810944   35         20000         0         1
2    15668575   26         43000         0         0
3    15603246   27         57000         0         0
4    15804002   19         76000         0         1
```

```
[29]: data.tail()
```

```
[29]:      User ID  Age  EstimatedSalary  Purchased  Gender_Male
395  15691863   46         41000         1         0
396  15706071   51         23000         1         1
397  15654296   50         20000         1         0
398  15755018   36         33000         0         1
399  15594041   49         36000         1         0
```

```
[30]: data.describe()
```

```
[30]:      User ID      Age  EstimatedSalary  Purchased  Gender_Male
count  4.000000e+02  400.000000         400.000000  400.000000  400.000000
mean    1.569154e+07  37.655000        69742.500000    0.357500    0.490000
std     7.165832e+04  10.482877        34096.960282    0.479864    0.500526
min     1.556669e+07  18.000000        15000.000000    0.000000    0.000000
25%     1.562676e+07  29.750000        43000.000000    0.000000    0.000000
50%     1.569434e+07  37.000000        70000.000000    0.000000    0.000000
75%     1.575036e+07  46.000000        88000.000000    1.000000    1.000000
max     1.581524e+07  60.000000       150000.000000    1.000000    1.000000
```


```
[31]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
↪4,random_state=100)
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model
```

[31]: LogisticRegression()

```
[32]: from sklearn.linear_model import Ridge
      clf=Ridge()
      clf.fit(x_train,y_train)
```

```
[32]: Ridge()
```

```
[100]: import seaborn as sns
      import matplotlib.pyplot as plt
      sns.lmplot(x="Age",y="EstimatedSalary",data=data,scatter_kws={"color":'green',
      ↪},line_kws={"color":'red'})
      sns.set_style('whitegrid')
      ax=plt.gca()
      plt.gca()
      plt.gca().set_facecolor('white')
```



```
[102]: data=pd.get_dummies(data,columns=["Gender_Male"],dtype=int,drop_first=True)
      data
```



```
[102]:
```

	User ID	Age	EstimatedSalary	Purchased	Gender_Male_1
0	15624510	19	19000	0	1
1	15810944	35	20000	0	1
2	15668575	26	43000	0	0
3	15603246	27	57000	0	0
4	15804002	19	76000	0	1
..	...	...	...	...	...
395	15691863	46	41000	1	0
396	15706071	51	23000	1	1
397	15654296	50	20000	1	0
398	15755018	36	33000	0	1
399	15594041	49	36000	1	0

[400 rows x 5 columns]

```
[104]: import numpy as np
```

```
[110]: #To check duplicate values
duplicate_rows=data.duplicated()
data[duplicate_rows].sum()
```

```
[110]: User ID      0
Age             0
EstimatedSalary 0
Purchased       0
Gender_Male_1   0
dtype: int64
```

```
[111]: print("Before dropping duplicate:",data.shape)
data.drop_duplicates()
print("After dropping duplicate:",data.shape)
```

Before dropping duplicate: (400, 5)


After dropping duplicate: (400, 5)

```
[115]: # Based on index value try to check the performance
response=data['Gender_Male_1']
response.dtype
```

```
[115]: dtype('int32')
```


```
[125]: sns.jointplot(x=response.index,y="Age",data=data,kind='hex')
```

```
[125]: <seaborn.axisgrid.JointGrid at 0x2145ce6c3e0>
```




```
[126]: sns.violinplot(response)
```

```
[126]: <Axes: ylabel='Gender_Male_1'>
```



```
[65]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lmplot(x="Age",y="EstimatedSalary",data=data,scatter_kws={"color":'yellow',
↪},line_kws={'color':"red"})
sns.set_style('whitegrid')
ax=plt.gca()
plt.gca()
plt.gca().set_facecolor('white')
```



[ ]: